
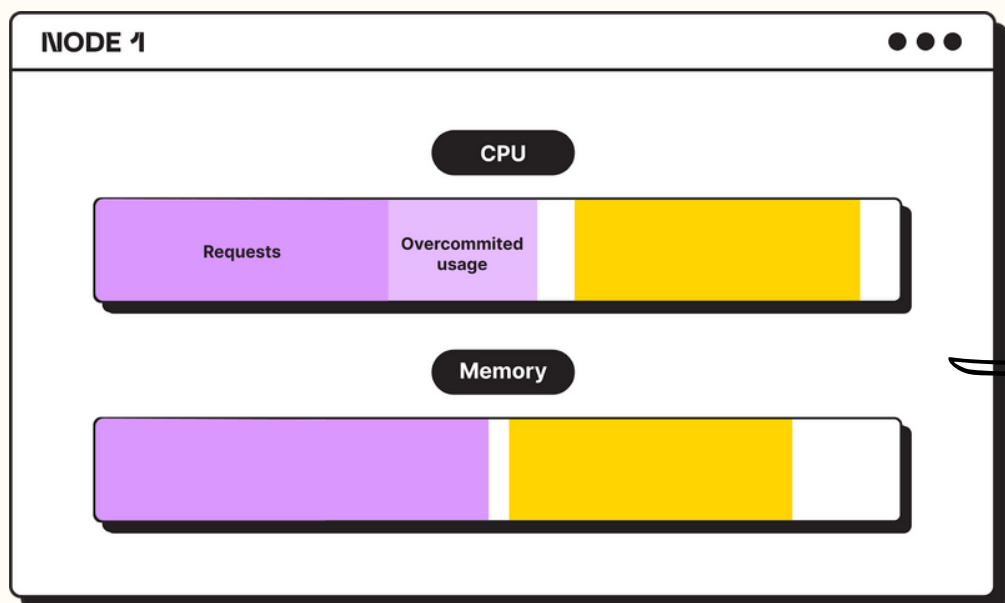


# Preventing OOMKilled Errors in Kubernetes




If you don't set the correct memory **limits and requests**, and/or if K8s doesn't assign proper QoS classes for some reason...





# You could end up with an OOMkilled error (or Out of Memory Killed error)


```
State:      Running
  Started:   Wed, 10 Mar 2021 11:35:31 +1100
Last State: Terminated
  Reason:    OOMKilled
Exit Code:  137
  Started:   Wed, 10 Mar 2021 08:06:03 +1100
Finished:    Wed, 10 Mar 2021 11:35:30 +1100
```



**OOMKilled errors occur when the Linux kernel's OOM Killer decides to terminate a process because it needs to free up some memory.**




**K8s should automatically move Pods to different nodes based on memory availability before it becomes necessary for one Pod's kernel to **start terminating** Pods due to low memory availability, but you still need to manage things right.**



The way to manage memory inside K8s via limits and requests in such a way that **OOMkilled errors** won't happen because memory will always be distributed properly between containers and Pods.



**Kubernetes can track the available **memory** on each node in a cluster and distribute Pods on different nodes accordingly.**




# Set the limit too high?

You might find yourself in a “noisy neighbors” situation when some applications affect the performance of others.





**Set the limit too low?**  
**And your application pods might**  
**crash spontaneously without**  
**you even knowing.**



As long as you set the **right memory limits** and requests – and assuming your cluster's total memory availability is sufficient to meet the needs of your workloads – you should be able to avoid OOMkilled errors, and still keep your deployments separated and stable.



# Thanks for reading!

---

If you find this useful, don't  
forget to save it!



## groundcover